

All prompts — the full comparison

Every prompt we tried for MedGemma-27B, and how it did. All numbers are **pooled over 1994 reports** (Zoe n=1495 + Maria n=499), **Core F1** and whole-report accuracy vs the reference annotator **LD**, with the second annotator **SG** held out. Baselines: the paper’s **Mistral-7B** (Tian et al.) and the **human** second annotator (SG vs LD).

The prompts

Prompt	One line	What it changes
v1	Baseline	Impression-first, short definitions — the starting point
v2	Professor’s revision	Neurologist role, read the report <i>body</i> first, extended ACNS/ILAE definitions
v3	Focal-epi exclusions	v1 + an explicit list of what does not count as focal epileptiform (raises precision)
v4	Procedure	v1 + a step-by-step detect → localize → assign routine
v5	Slowing discriminator	v3 + a focal-vs-generalized rule for non-epileptiform slowing
v6	Ask for consistency	v3 + a prompt instruction to self-check the overall/subtype agreement
v7	Body-aware abnormality	v5 + “call it abnormal if the body describes an abnormality the Impression downplays”
v8	Simplified	A deliberately short prompt (~340 tokens) — only the two decision boundaries
v9	Reasoning-first	Concise expert frame + a free-text reasoning field, letting the model deliberate
v10	Evidence calibration	v5 + “tie confidence to how explicitly the report supports each subtype”

A **...g** suffix (e.g. **v5g**) means the variant is run with **grammar-enforced consistency** (ENFORCE_CONSISTENCY=1): a GBNF grammar emits `overall_abnormal` last and only lets it be abnormal when a subtype is, so a self-contradictory answer is impossible to produce. Full prompt texts are in [./prompts/](#).

Where every prompt lands

What you see: whole-report accuracy (all five labels correct) for each configuration, from Mistral-7B up to the human agreement line (89.8%, dotted). Our baseline **v1 already clears Mistral by ~10 points**. Adding grammar-enforced consistency is the big step (**v3g/v5g**), and **v5g (87.6%)** is the best — within ~2 points of the human ceiling. The rejected ideas cluster below it: **v8 (simplified)** even falls *below* the v1 baseline, because dropping the detailed rules brings the over-calling straight back.

Our top prompts vs Mistral and human, per category

What you see: Core F1 for **each label scored on its own**, for our three strongest prompts (**v5g, v3g, v7g**), Mistral-7B, and the human annotator. Note these are a *different metric* from the first chart: the first chart’s 87.6% is the share of reports with **all five** labels correct at once; here each number

MedGemma-27B — whole-report accuracy by prompt · pooled n=1994

each dot = one configuration · dotted line = human annotator agreement (89.8%)




Figure 1: Whole-report accuracy by prompt

is the F1 of **one** label — so the two are not meant to match. (Grammar variants shown at **Q4_K_S** — the same quant as the Full-numbers table, so the dots match those rows exactly.) Our prompts (blue / orange / violet) **cluster tightly together and sit far above Mistral (grey)** on the harder three classes (Gen Epi, Focal Non-epi, Gen Non-epi) — the gap to Mistral is 13-14 points there. Against the **human (green)** the cluster sits right at the human level: **at or slightly above it on the epileptiform classes** (Focal Epi ~86 vs 86, Gen Epi ~88-89 vs 88) and on **Gen Non-epi** (~89-90 vs 90), trailing only on **Abnormality** and **Focal Non-epi** — where even the two humans disagree most. That the three prompts land so close to each other is itself the point: the result is stable across our top variants, not a fluke of one prompt.

Why we win — confidence, not just direction

What you see: each line runs from **Core F1** (● — did we get present/absent right?) to **Certainty F1** (○ — did we get the exact 1-4 confidence level?); the line length is how much is lost when the exact level is required. Our **v5g (blue)** beats Mistral on core almost everywhere, but the real separation is the **exact level**: Mistral's open circles collapse (Focal Epi 83 → 41, Focal Non-epi 76 → 45, Gen Non-epi 75 → 52) — it points in the right direction but badly misjudges *how sure* to be — while v5g holds far better (Focal Epi 86 → 67, Focal Non-epi 89 → 63). Getting the confidence level right, not just the yes/no, is where our prompt+grammar approach pulls ahead.

Which quantization, and does it generalize?

What you see: v5g on the two model sizes. They tie on whole-report accuracy (87.6% each); per category **Q2_K** is stronger on the epileptiform classes (Focal Epi 88 vs 86) while **Q4_K_S** is stronger on the diffuse/slowing classes (Gen Non-epi 89 vs 87). Pick Q2_K for half the footprint, Q4_K_S if encephalopathy/diffuse findings matter most.

What you see: the same model on the **seen** neurologist (Zoe) and an **unseen** one (Maria). Performance holds across reporting styles — the epileptiform and focal classes are as strong or stronger on the unseen data — so the model generalizes rather than fitting one annotator's phrasing.

Per-category F1 — our top prompts vs Mistral-7B vs human

each label scored on its own (Core F1, 0-100) · NOT the whole-report % · grammar variants at Q4_K_S · pooled n=1994



Figure 2: Per-category F1 — top prompts vs Mistral and human




Figure 3: v5g vs Mistral — core vs certainty




Figure 4: Q2 vs Q4 per category

Generalization — v5g on the seen vs an unseen neurologist




Figure 5: Generalization — Zoe vs Maria

How the metrics work (read before the tables)

Every label uses the **1-4 scale** (1 confident-no · 2 low-no · 3 low-yes · 4 confident-yes). We score with **F1 throughout**, at two strictness levels:

- **Core F1** — collapse to yes/no (1-2 = absent, 3-4 = present) and score how well “present” is detected. “3 vs 4” counts as correct (both mean present).
- **Certainty F1** — the exact 1-4 must match; “3 vs 4” is wrong (right direction, wrong confidence). The harder test — it shows whether the model captures *how sure* to be.

Why F1, not plain accuracy? The classes are very imbalanced (Focal/Gen Epi are present in only ~5% of reports). Plain accuracy would score ~98% just by always predicting “absent”, hiding the failure to catch the rare findings. $F1 = 2 \cdot TP / (2 \cdot TP + FP + FN)$ ignores the easy true-negatives and measures only how well the *present* class is actually caught (balancing precision and recall). It is the paper’s core metric.

The one exception is Whole-report (All-5) — the share of reports where *all five* labels are correct. That is an exact-match count by nature (there is no F1 for a whole report), so it stays a plain percentage.

Full numbers (pooled n=1994)

Per-category values are **Core F1**; “Whole” is **All-5 accuracy** (share of reports with all five labels correct), vs LD and vs the held-out annotator SG.

Prompt	Abnorm	Focal Epi	Gen Epi	Focal Non	Gen Non	Whole vs LD	Whole vs SG
Mistral-7B	94.7	82.8	74.8	75.6	75.2	74.5	—
v1	96.8	80.4	88.5	88.0	84.0	83.9	82.9
v3	95.6	82.7	87.0	87.5	85.7	84.2	82.4
v3g	95.9	86.4	89.2	86.9	89.6	86.5	86.0
v5g Q2	95.4	87.6	89.9	88.8	86.7	87.6	85.6
v5g Q4	96.0	85.6	88.8	89.0	89.3	87.6	86.2
v7g	96.2	86.0	87.8	87.4	90.0	87.1	86.1

Prompt	Abnorm	Focal Epi	Gen Epi	Focal Non	Gen Non	Whole vs LD	Whole vs SG
v8g	93.9	76.9	85.4	85.8	87.7	83.6	82.8
v10g	96.1	86.9	87.8	87.4	89.8	86.7	86.1
Human SG	98.0	85.7	87.5	90.8	90.0	89.8	—

(v9g omitted from the pooled table — it is much slower and did not finish the full run; on a fair same-case comparison it was 3–6 points **below** v5g. Values shown use Q4 for the grammar variants except where a quant is named.)

Certainty (exact-level) F1

The Full-numbers table above is **Core F1**. Here is the same, per category, at the stricter **Certainty** level — the exact 1–4 confidence must match. This is the harder test, and it is where our approach separates most from Mistral.

Model	Abnorm	Focal Epi	Gen Epi	Focal Non	Gen Non
Mistral-7B	72.4	41.4	69.2	44.8	52.2
v1	56.0	67.6	78.1	45.2	54.7
v3	59.8	69.2	74.6	50.5	56.5
v3g (Q4)	75.6	66.3	75.9	60.6	67.1
v5g (Q2)	65.8	76.8	83.6	59.3	60.1
v5g (Q4)	74.1	66.7	72.4	63.1	69.0
v7g (Q4)	77.6	66.0	74.5	62.2	69.6
v8g (Q4)	61.6	59.4	73.4	50.4	61.1
v10g (Q4)	75.1	67.7	77.6	62.0	71.5
Human (SG)	92.0	72.0	84.1	46.9	58.3

The rare **epileptiform** classes are the story: at the exact level Mistral collapses (Focal Epi 41) while our grammar prompts hold (66–77). On Abnormality the human is far ahead (92) — judging *how sure* to be about “abnormal” is the hardest thing to imitate. On the slowing classes even the human’s certainty F1 is modest (Focal Non 47, Gen Non 58), and our models sit at or above it.

Regenerate: `python -m analysis.fl_tables.`

What worked, what didn’t

- **Won:** the combination **v5 prompt + grammar-enforced consistency (v5g)**. The grammar is the single biggest lever — it removes every overall/subtype contradiction *and*, by letting the model decide the parts before the whole, lifts the epileptiform F1.
- **Rejected — asking instead of enforcing (v6):** the model doesn’t hold the consistency rule on its own.
- **Rejected — body-aware abnormality (v7):** the remaining missed-abnormal cases are genuinely ambiguous, not a prompt gap.
- **Rejected — simplification (v8):** the detailed focal-epileptiform exclusions are load-bearing; removing them re-introduces over-calling.
- **Rejected — reasoning-first (v9):** given latitude to deliberate, the model’s unaided judgement still under-performs the explicit rules by 3–6 points.
- **Rejected — evidence calibration (v10):** ~1 point below v5g on both quants; the “prefer absent” rule slightly hurt the rare Focal Epi class on Q2.

The improvement **holds against the held-out annotator SG** (82.9% → 86.2%), so it is a real gain, not fitting to LD; the residual gap to the human is close to the level of human-human disagreement, which is a reason to validate carefully rather than a claim that nothing more is possible.

Per-prompt detail — core vs certainty, Q2 vs Q4

One chart per prompt, each overlaying both quantizations. The **filled dot** is Core F1 (present/absent), the **open circle** is strict Certainty F1 (exact 1-4 level); the line length is the drop when the exact level is required. All pooled over 1994 reports.

v5g (best):




Figure 6: v5g core vs certainty

v3g:

v7g:

v8g (simplified):

v10g (calibrated):

Mistral-7B (reference):

Reproduce

```
pip install -e .
python -m analysis.make_tables           # the numbers tables
python -m analysis.plot_all_prompts     # every chart in this report -> reports/figures/
python -m analysis.validate_vs_sg       # cross-annotator validation
```

Per-run job IDs and hypotheses: `../experiments/`. Model: MedGemma-27B GGUF (Q2_K / Q4_K_S), llama.cpp grammar-constrained, temperature 0, 64-core CPU.

MedGemma-27B · v3g — core vs certainty (Q2 vs Q4)

● core (present/absent) ○ exact level 1-4 · pooled n=1994




Figure 7: v3g core vs certainty

MedGemma-27B · v7g — core vs certainty (Q2 vs Q4)

● core (present/absent) ○ exact level 1-4 · pooled n=1994




Figure 8: v7g core vs certainty

MedGemma-27B · v8g (simplified) — core vs certainty (Q2 vs Q4)

● core (present/absent) ○ exact level 1-4 · pooled n=1994




Figure 9: v8g core vs certainty

MedGemma-27B · v10g (calibrated) — core vs certainty (Q2 vs Q4)

● core (present/absent) ○ exact level 1-4 · pooled n=1994




Figure 10: v10g core vs certainty

Mistral-7B — core vs certainty

F1: binary present/absent vs strict exact-level match · pooled n=1994




Figure 11: Mistral core vs certainty