

Harvard EEG Database — ID structure of the label tables

The HEEDB label tables (*_EEG_with_reports, one per hospital) are organized as a simple hierarchy of identifiers. From the top:

```
SiteID           hospital           S0001 = MGH, S0002 = BWH, I0002 = BIDMC, I0003 = BCH
└─ BDSPPatientID patient           de-identified; the master key
    └─ report      NeuroReport_<a>_<b>_<date>.txt  ( = DeidentifiedName(Reports) )
        └─ SessionID recording segment under that report
```

SiteID — the hospital. There is one label table per site; SiteID is constant within a table.

BDSPPatientID — the de-identified patient, and the key that ties everything together. One patient has many EEG reports over time. Patient-level facts (diagnoses in ICD-10, medications in ATC) are attached here, at the patient level — not per report.

Report — DeidentifiedName(Reports), a filename of the form NeuroReport_<a>__<date>.txt. This is one complete EEG report document (Impression / Clinical indication / Method / Description). The free-text of the report is **not stored in the database** — it lives in the HEEDB zip dumps and is joined back by this exact filename.

SessionID — a small segment counter *within* a report. A continuous / long-term monitoring study (EMU or ICU cEEG) is split into many recording segments, so one report covers many SessionIDs. SessionID is not unique and carries no separate text — every session of a report shares the same report document.

Because of this, the table has more rows than reports, and more reports than patients. For MGH:

Level	Count (MGH)
rows (session segments)	90,716
unique reports (DeidentifiedName(Reports))	55,557
unique patients (BDSPPatientID)	34,620

So the ordering is: **one patient → several reports over time → each report → several recording sessions**; diagnoses and medications hang off the patient, and the report text is fetched from the zip dumps by the report filename.